

# 15CSE334 -Big Data Analytics

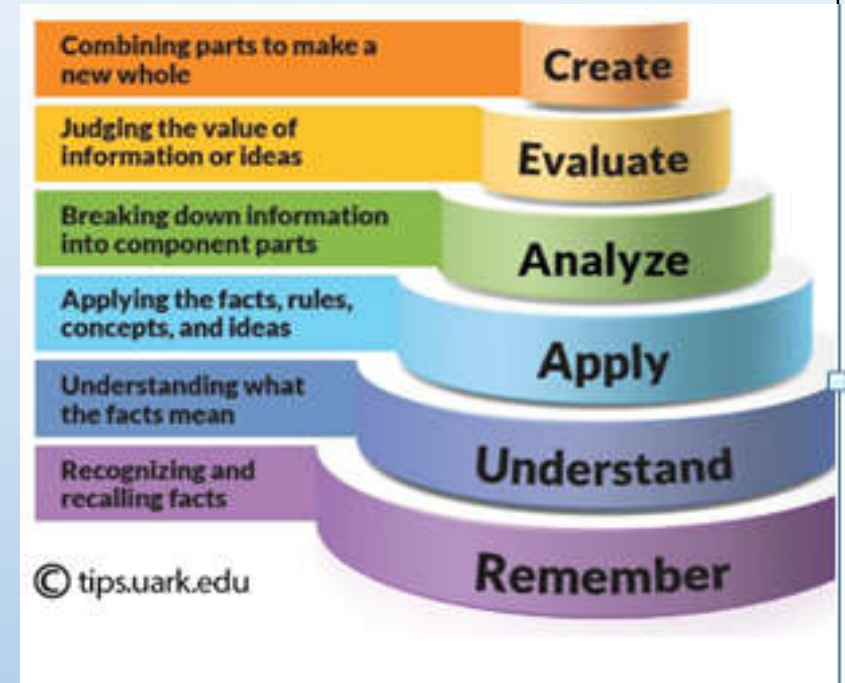
Dr. R. Karthi

# Course Outline

- Understand the basics of Big data.
- Understand Hadoop architecture and tools used for big data.
- Learn to implement map-reduce programs for big data.
- Understand and apply Big data tools for various applications.
- Course – Lab Based Course ( 3 Credits)
  - 3 hrs. theory + 3 hrs. lab per week
  - Laptop with internet connectivity required for Theory and Lab hrs.
  - Enthusiasm to explore and learn new skills.
  - Theory Exams will be in paper based mode.
  - Lab slots are used to demonstrate and have an hand on experience on each platform/tool.

# Course Outcomes

	Course Outcome	Bloom's Taxonomy Level
CO 1	<b>Understand</b> fundamental concepts of Big Data and its technologies	L2
CO 2	<b>Apply</b> concepts of Map Reduce framework for optimization.	L3
CO 3	<b>Analyze</b> appropriate No SQL database techniques for storing and processing large volumes of structured and unstructured data.	L4
CO 4	<b>Apply</b> data analytics solutions using Hadoop ecosystems	L3
CO 5	<b>Explore</b> tools for machine learning and reporting	L4



# Evaluation Pattern

Evaluation Component	Weightage
Periodical 1	15
Periodical 2	15
Continuous assessment (Quiz/Tutorial/Assignment)	20
End Semester	50

# Course Mapping

- Introduction to Big Data
- Introduction to Hadoop
  - Architecture of hadoop
  - Map Reduce programming
- Hadoop ecosystem –tools
  - Pig
  - Hive
- NOSQL Databases
  - MongoDB
  - Cassandra
- Machine Learning

## Points to Remember

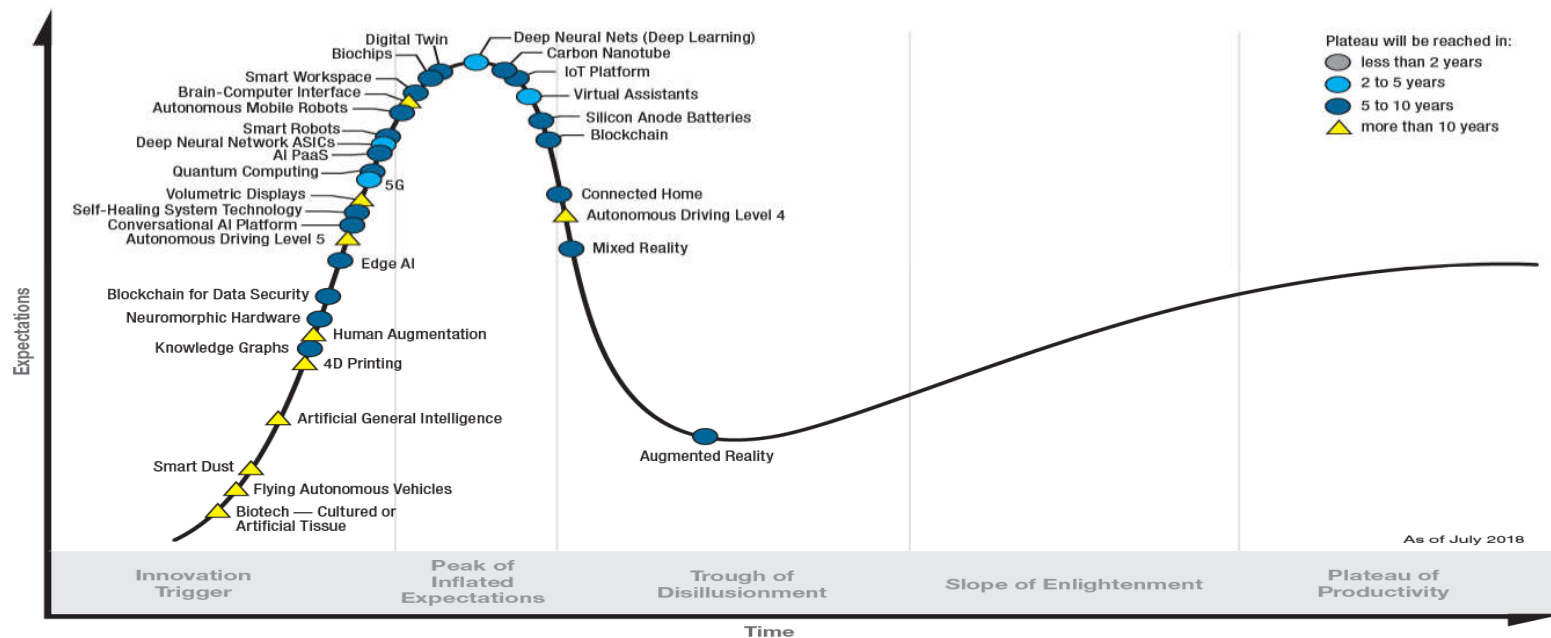
- Maintain a Observation notebook for lab.
- All exercises and evaluations done in lab and class are to be recorded.
- Keep your laptops updated with sufficient processing speed and memory.
- Amrita AUMS / Email accounts to be active.
- Case study on big data for real world applications.

# Introduction to Big Data Analytics

Dr. R. Karthi

# Emerge in the Gartner Hype Cycle for Emerging Technologies, 2018

## Hype Cycle for Emerging Technologies, 2018



[gartner.com/SmarterWithGartner](https://gartner.com/SmarterWithGartner)

Source: Gartner (August 2018)  
© 2018 Gartner, Inc. and/or its affiliates. All rights reserved.

**Gartner**



# Emergence in the Gartner Hype Cycle for Emerging Technologies, 2019



# Big Data – Industry Needs

A 2011 study by the McKinsey Global Institute predicts that by 2018-19 the **U.S. alone will face a shortage of 140,000 to 190,000 people with deep analytical skills** as well as ***1.5 million managers and analysts*** to analyze big data and make decisions.

# What Is Big Data?

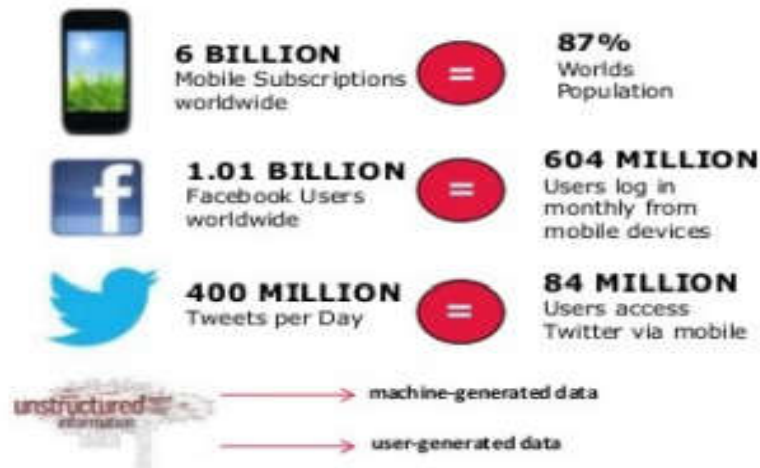
- ❑ No single standard definition...

- ❑ *Extremely large data sets that may be analyzed computationally to reveal patterns, trends, and associations, especially relating to human behavior and interactions.*
- ❑ *Big data is an evolving term that describes any voluminous amount of structured, semi-structured and unstructured data that has the potential to be mined for information.*
- ❑ *“**Big Data**” is data whose scale, diversity, and complexity require new architecture, techniques, algorithms, and analytics to manage it and extract value and hidden knowledge from it...*

# What makes Data "Big"

## What Makes Big Data So Big?

pactera



1 Bit = Binary Digit

8 Bits = 1 Byte

1024 Bytes = 1 Kilobyte

1024 Kilobytes = 1 Megabyte

1024 Megabytes = 1 Gigabyte

1024 Gigabytes = 1 Terabyte

1024 Terabytes = 1 Petabyte

1024 Petabytes = 1 Exabyte

1024 Exabytes = 1 Zettabyte

1024 Zettabytes = 1 Yottabyte

1024 Yottabytes = 1 Brontobyte

1024 Brontobytes = 1 Geopbyte


**Big Data will get only bigger as traffic from smartphones and tablets outpaces traditional devices.**

Percentage of Web Traffic by 2016:



© Pactera. Confidential. All Rights Reserved.

Volume of Digital Content:



Source: International Data Corporation (IDC) "Worldwide Personal Datacenter Market: Forecast and Analysis, 2012-2016" (2012).  
Source: International Data Corporation (IDC) "Worldwide Personal Datacenter Market: Forecast and Analysis, 2012-2016" (2012).

# Who's Generating Big Data



**Social media and networks**  
(all of us are generating data)



**Scientific instruments**  
(collecting all sorts of data)



**Mobile devices**  
(tracking all objects all the time)



**Sensor technology and networks**  
(measuring all kinds of data)

- The progress and innovation is no longer hindered by the ability to collect data
- But, by the ability to manage, analyze, summarize, visualize, and discover knowledge from the collected data in a timely manner and in a scalable fashion



# The Model Has Changed...

## ➤ The Model of Generating/Consuming Data has Changed

**Old Model:** Few companies are generating data, all others are consuming data



**New Model:** all of us are generating data, and all of us are consuming data




## Classification of Digital Data

Digital data is classified into the following categories:

- Structured data
- Semi-structured data
- Unstructured data

Approximate percentage distribution of digital data

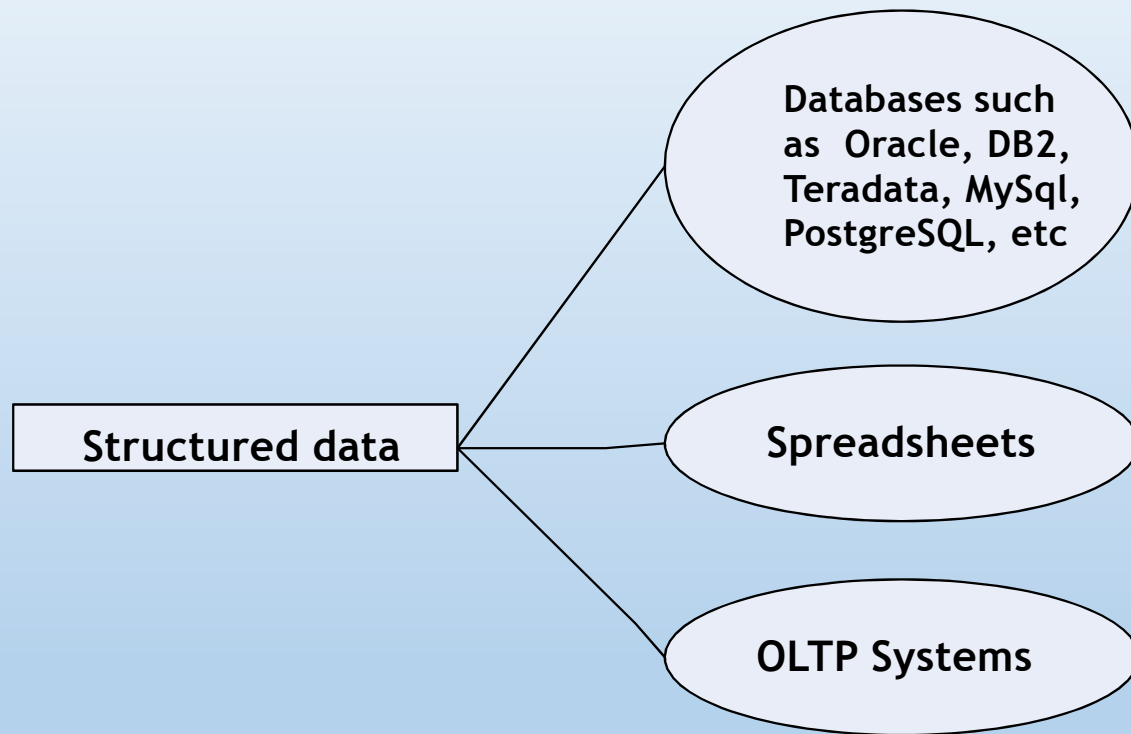


## Structured Data

- This is the data which is in an organized form (e.g., in rows and columns) and can be easily used by a computer program - Relational data model
- Specific Data type and Constraints(Unique, Not Null) are defined
- Relationships exist between entities of data, such as classes and their objects.
- Data stored in databases is an example of structured data.
  -
- Example : Employee Data base

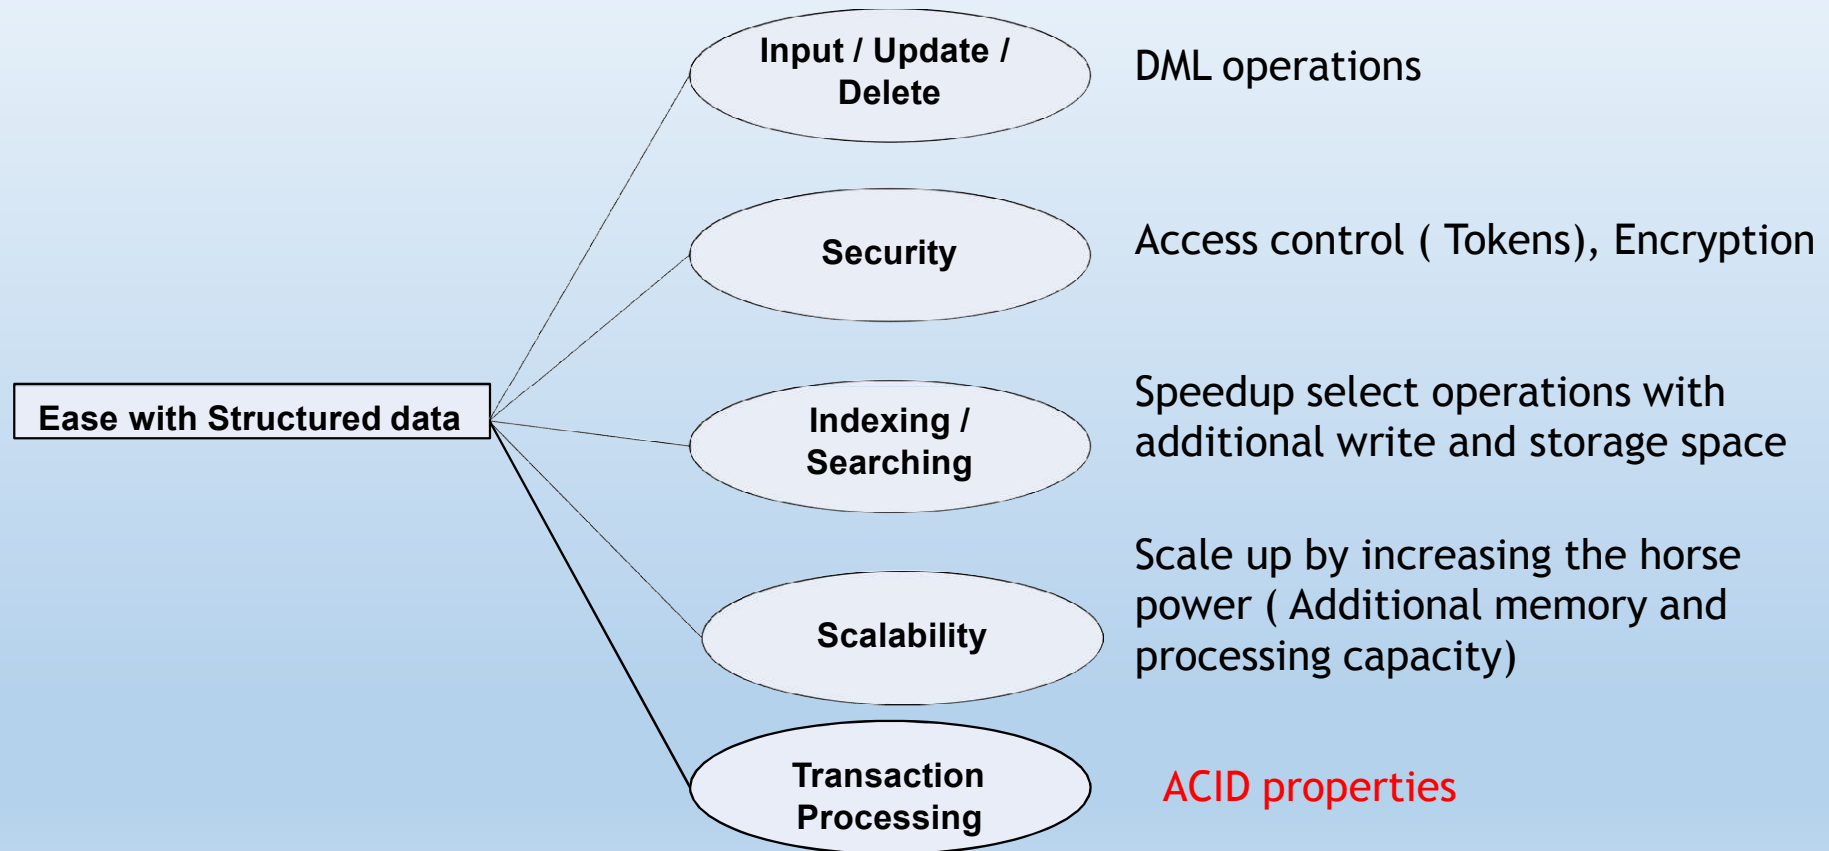


## Sources of Structured Data



Online Transaction Processing Systems

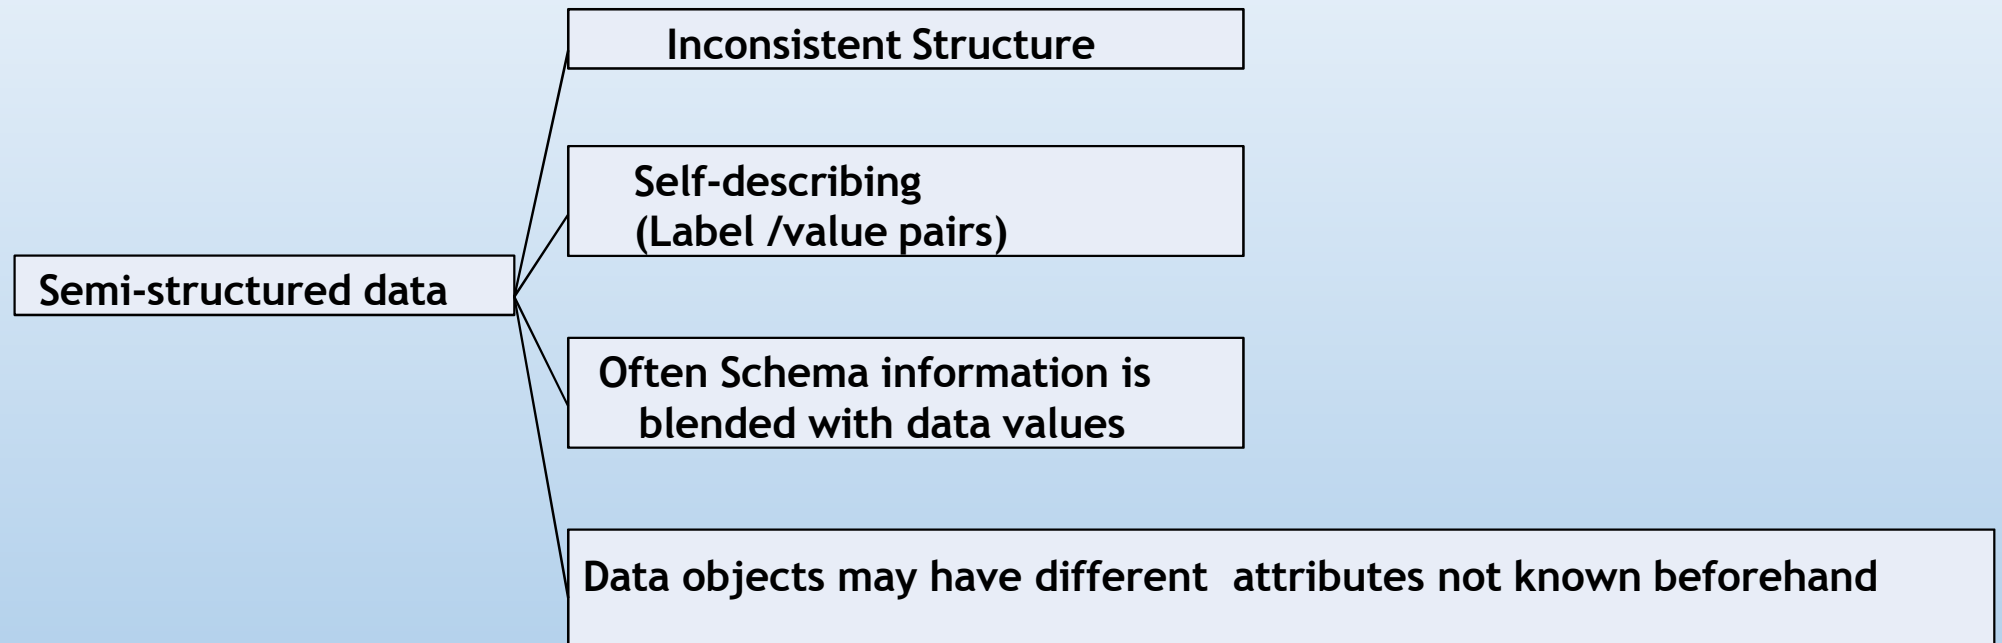
## Ease with Structured Data



## Semi-structured Data

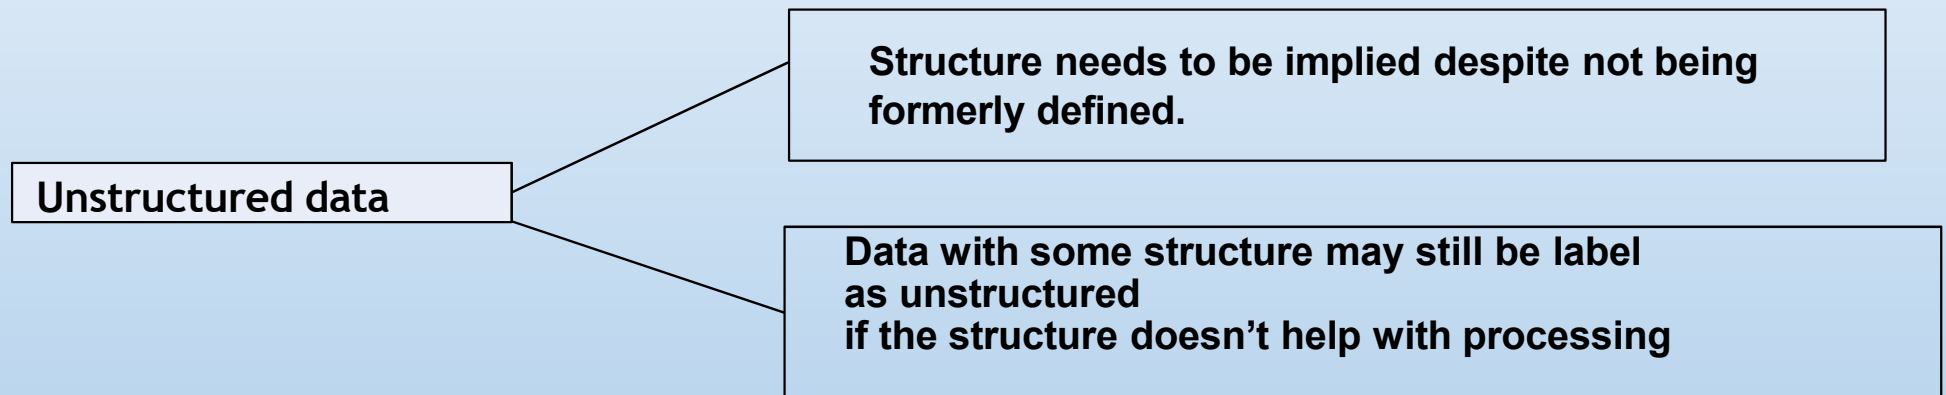
- This is the data which does not conform to a data model but has some structure. However, it is not in a form which can be used easily by a computer program.
- It uses tags to separate semantic elements and to enforce hierarchies of records and fields within data.
- No separation between schema and data.
- Metadata for this data is available but is not sufficient.
- Example: emails, XML, markup languages like HTML, etc.

## Characteristics of Semi-structured Data

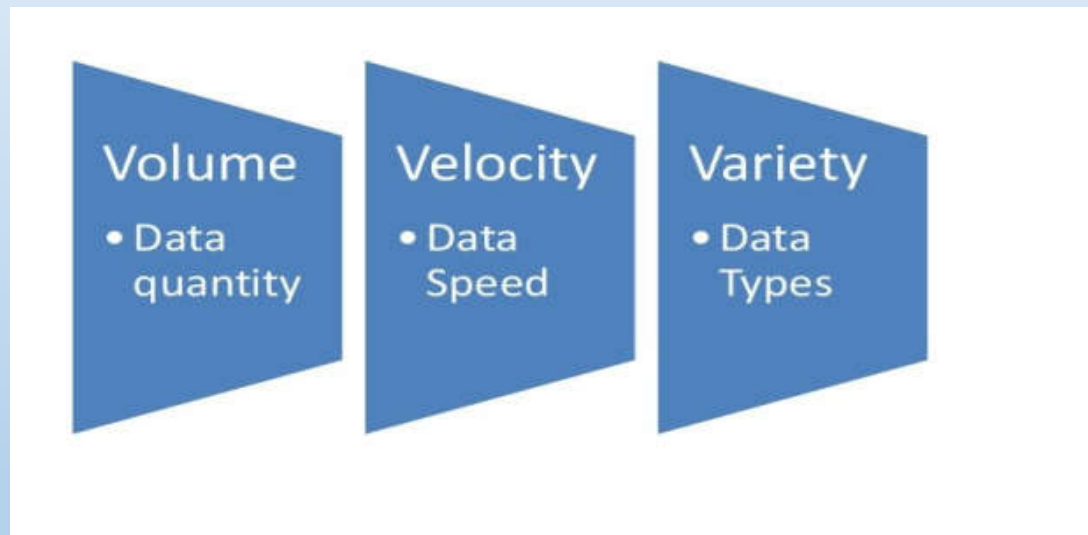


## Unstructured Data

- Data which does not conform to a data model or is not in a form which can be used easily by a computer program.
- About 80-90% data of an organization is in this format.
- Example: memos, chats , PowerPoint presentations, images, videos, letters, researches, white papers, body of an email, etc.

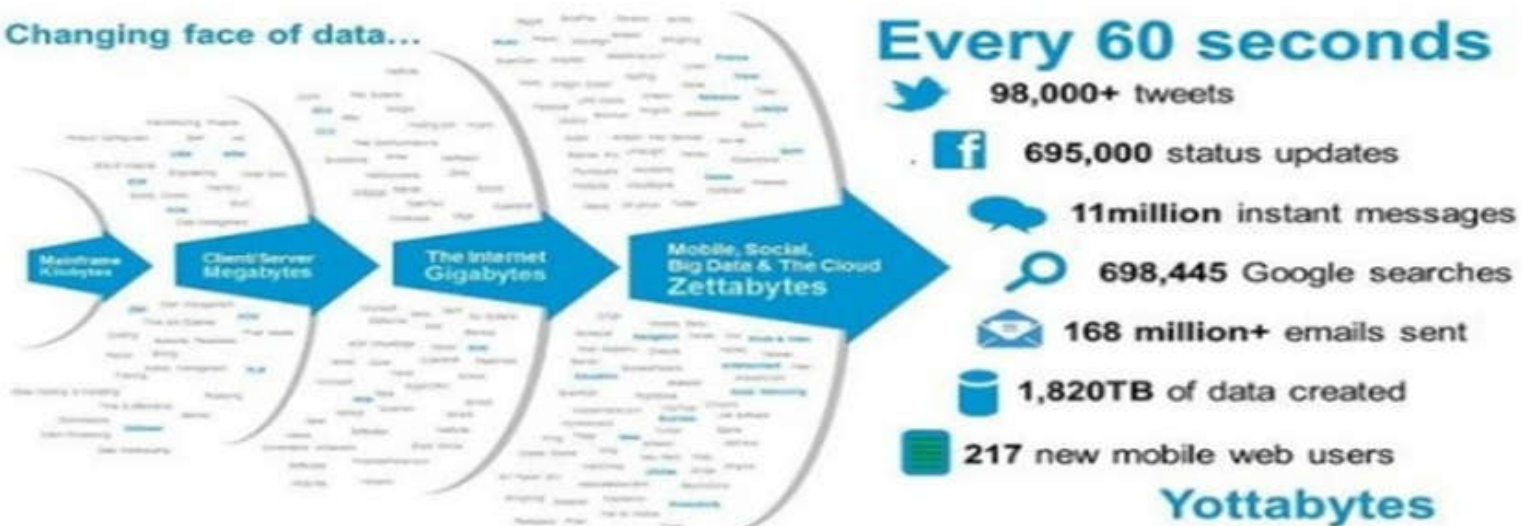


# Characteristics of Big Data

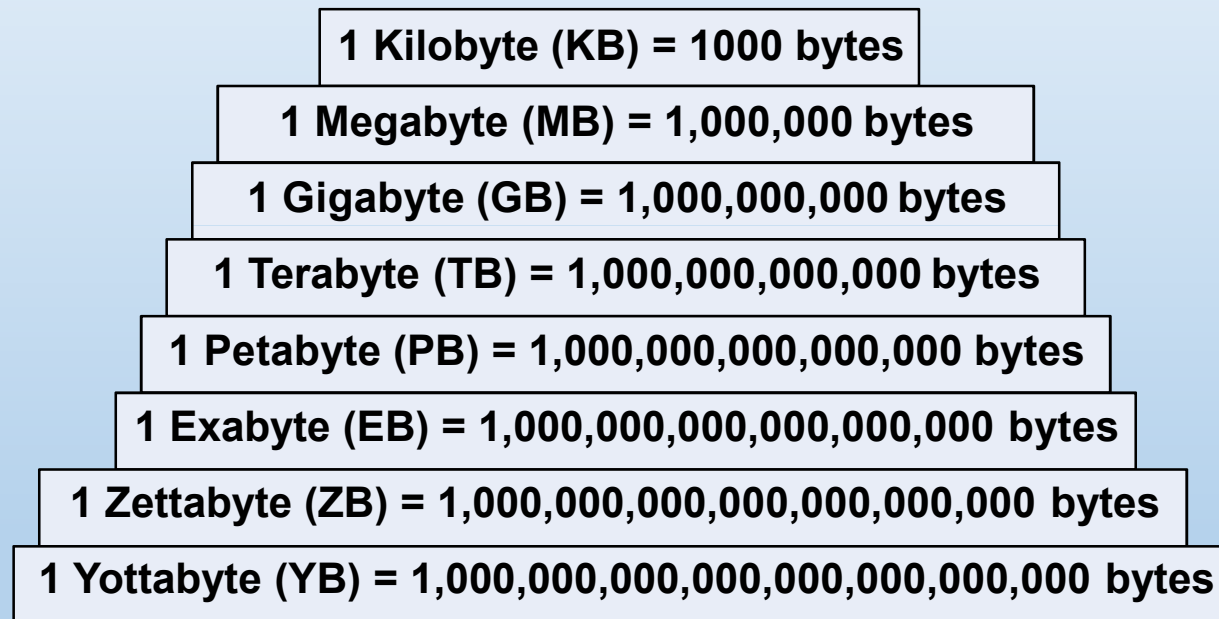


# Data Volume

## Data volume is increasing exponentially



# Volume





## Velocity

**Batch → Periodic → Near real time → Real-time processing**

# Characteristics of Big Data: Velocity

- Data is begin generated fast and need to be processed fast
- Online Data Analytics
- Late decisions → missing opportunities



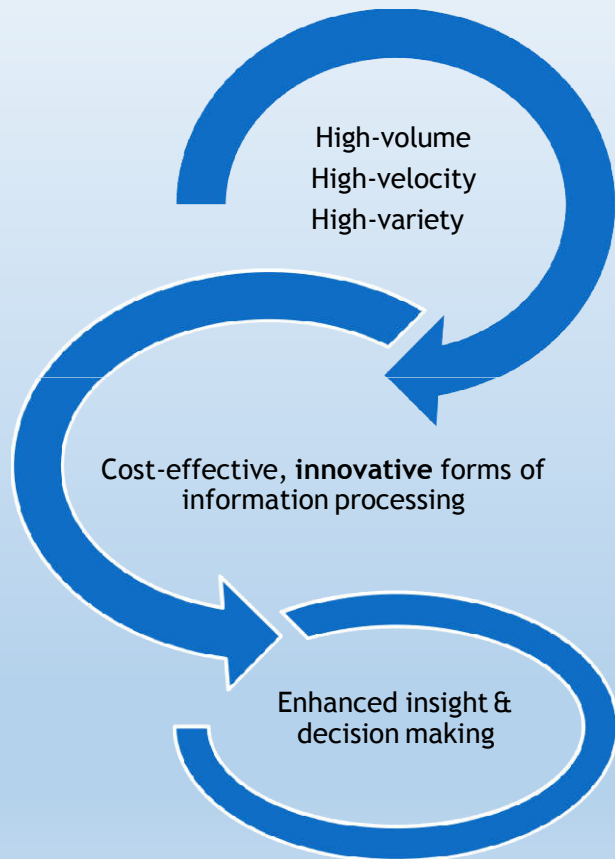
- **Examples**

- **E-Promotions:** Based on your current location, your purchase history, what you like → send promotions right now for store next to you
- **Healthcare monitoring:** sensors monitoring your activities and body → any abnormal measurements require immediate reaction

# Variety

- **Structured data:** example: traditional transaction processing systems and RDBMS, etc.
- **Semi-structured data:** example: Hyper Text Markup Language (HTML), eXtensible Markup Language (XML).
- **Unstructured data:** example: unstructured text documents, audio, video, email, photos, PDFs, social media, etc.

# Definition of Big Data



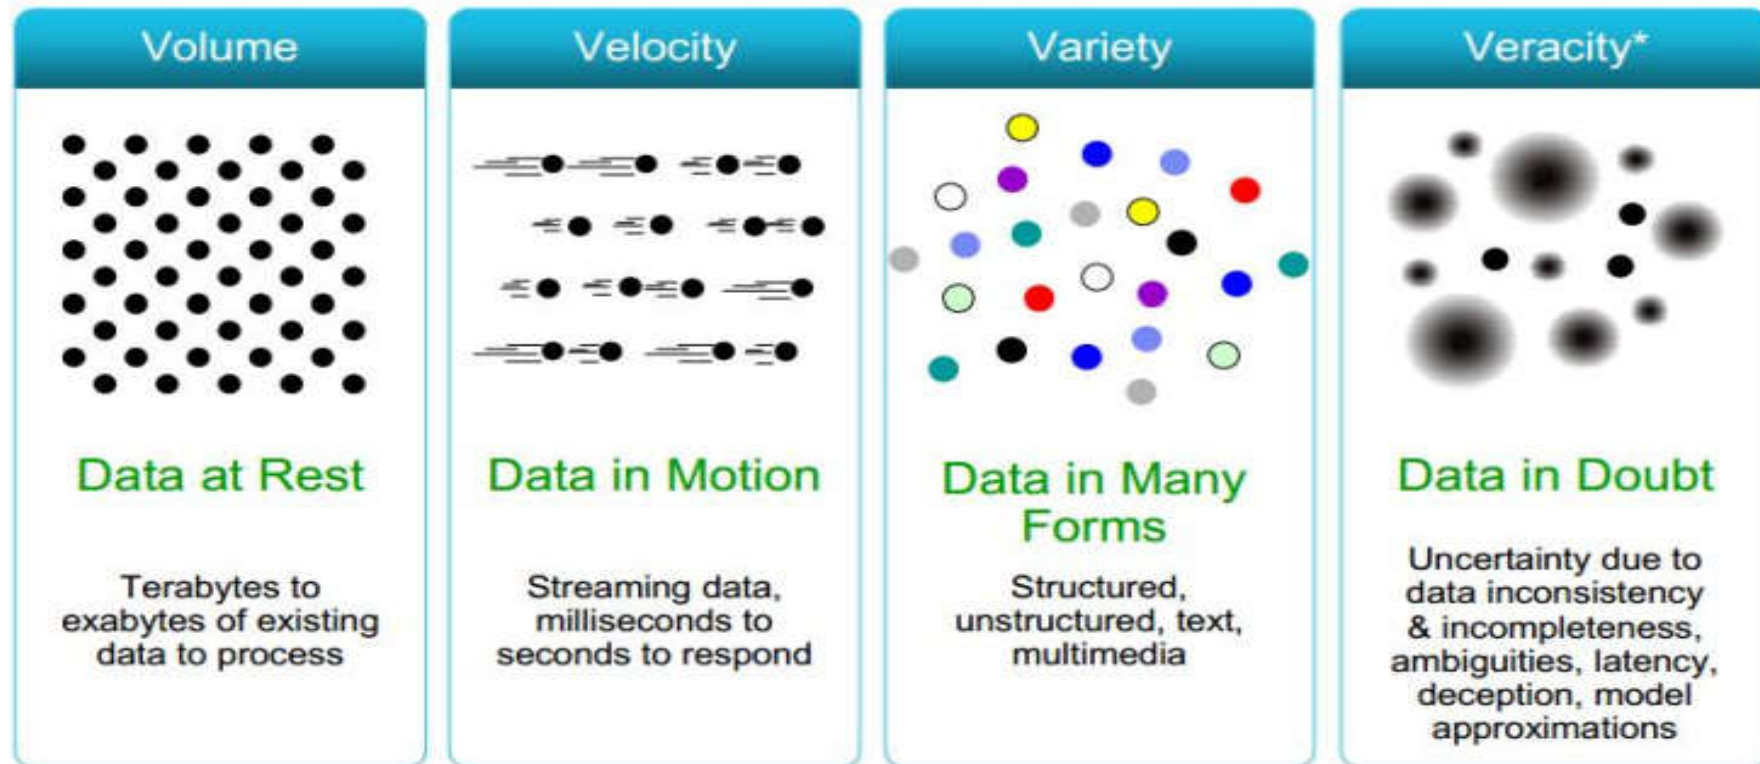
*Big Data is high-volume, high-velocity, and high-variety information assets that demand cost effective, innovative forms of information processing for enhanced insight and decision making.*

Source: Gartner IT Glossary


## Other Characteristics of Data - Traits of Big Data

- Veracity - the quality of being true or accurate.
- Volatility / Validity - how long is the data valid and how long should it be stored.
- Variability – Data whose meaning is changing , inconsistencies in the data.

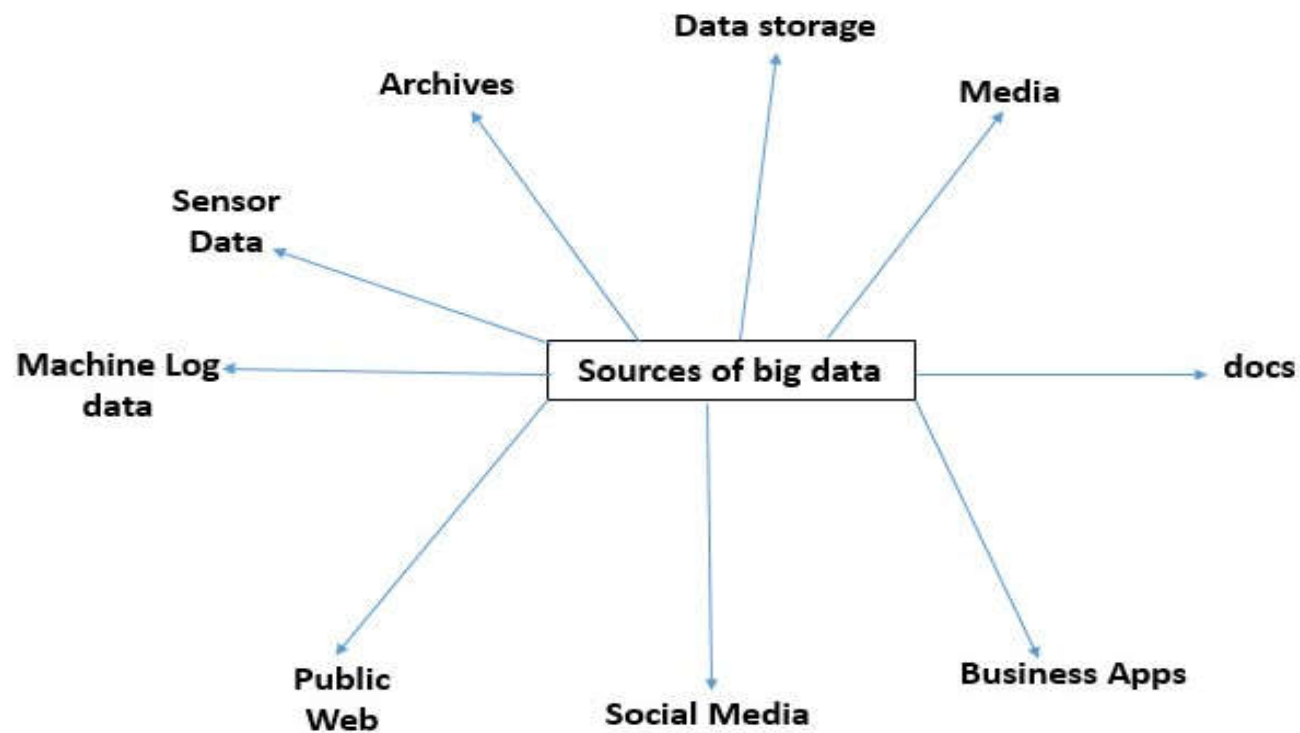
# Some Make it 4V's



# Challenges with Big Data



# Sources of Big Data





# Classification of analytics

Two thoughts:

1. Classify analytics into basic, operational, advanced and monetized.
2. Classify analytics into analytics1.0, 2.0 and 3.0

## **Basic analytics:**

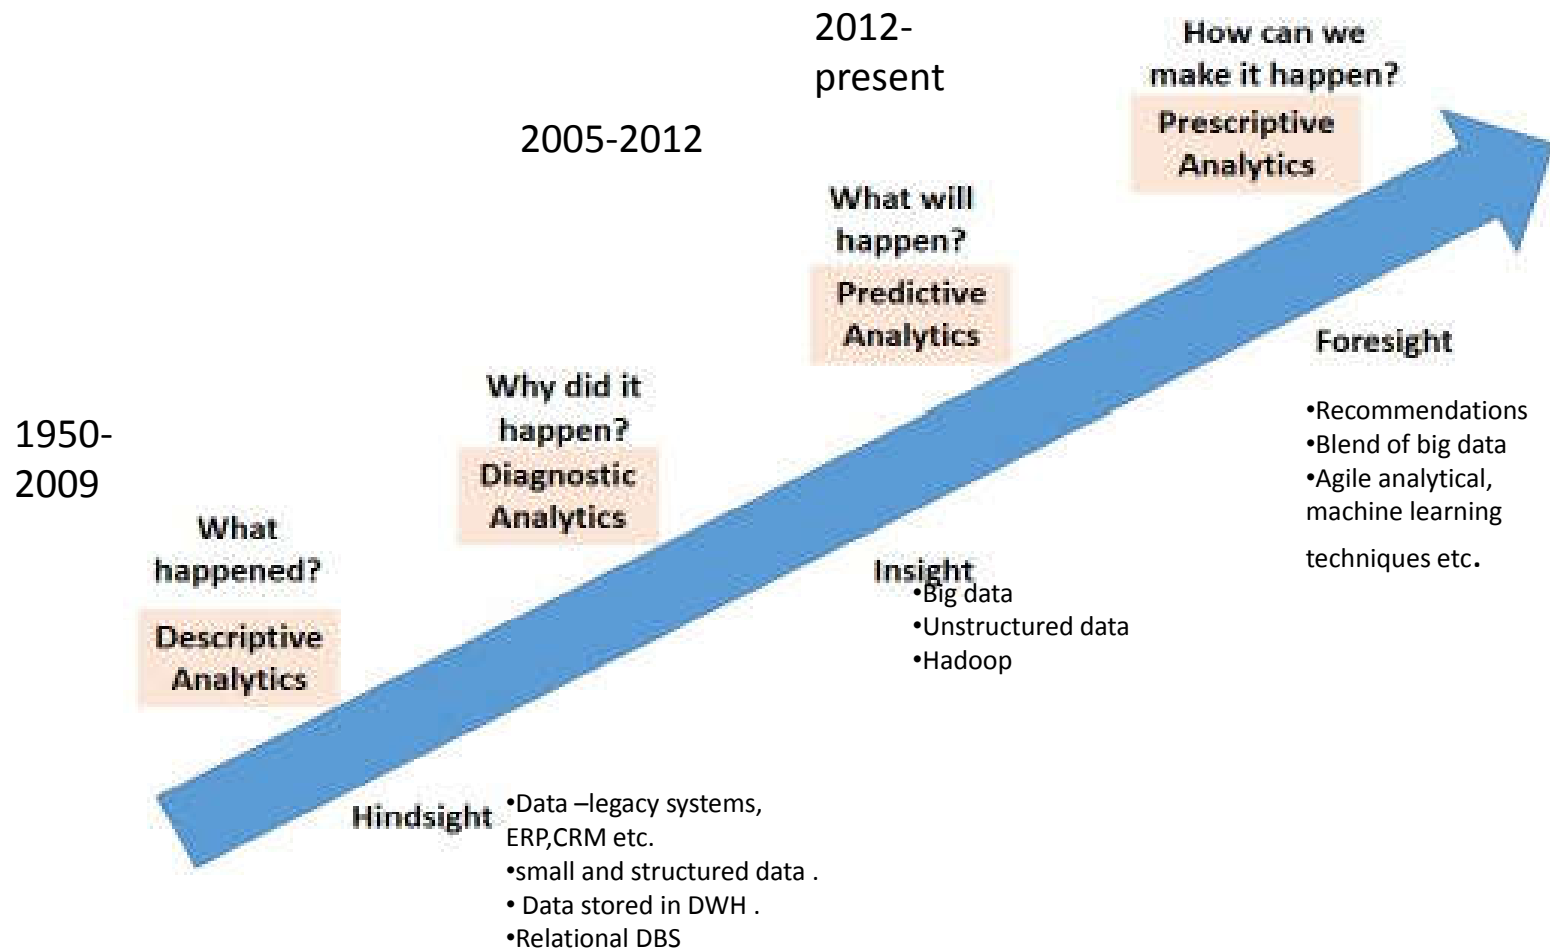
- Slicing and dicing data to help with basic business insights
- Based on historical data , visualization etc.

**Operational Analytics:** Enterprise business analysis (OLTP data)

**Advanced Prediction analytics:** Forecasting for the future/Prediction

**Monetized Analytics:** Analysis of data to derive direct business revenue.

## Analytics 1.0, 2.0 and 3.0



# Types of Analytics

- **Descriptive analytics:**

- What happened / why it has occurred.
- Reporting tools, OLAP, dash boards, data visualization.

- **Predictive analytics :**

- Model forecast an outcome at some future state or time based upon changes to the model inputs.
- Algorithms are regression analysis, time series regression, machine learning methods – CART and neural networks.


- **Prescriptive analytics:**

- Uses predictive models to suggest actions to take for optimal outcomes.
- Prescriptive analytics relies on optimization and rules-based techniques


Forecasting the load on the electric grid over the next 24 hours is an example of predictive analytics, whereas deciding how to operate power plant based on this forecast represents prescriptive analytics.

# Traditional Business Intelligence (BI) – (Data ware house) versus Big Data


# A Typical Data Warehouse Environment



# A Typical Hadoop Environment



# Co-existence of Big Data and Data Warehouse

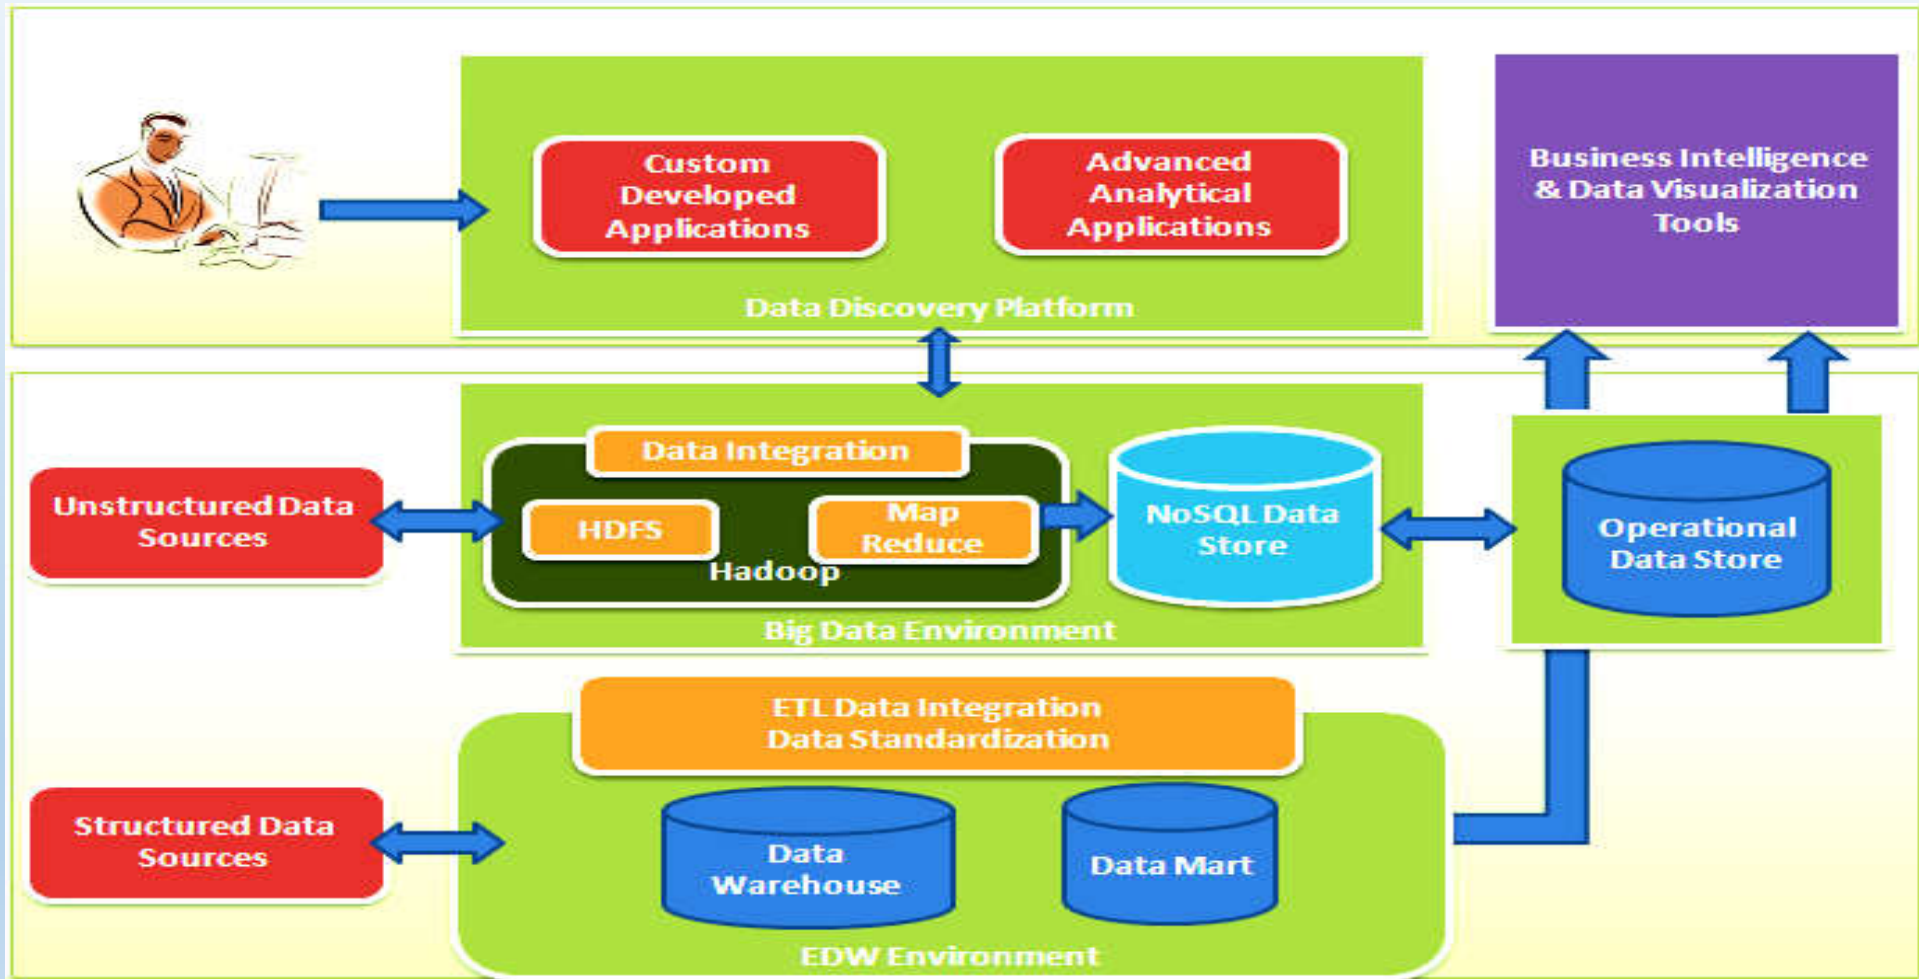


# Difference between Traditional BI and Big Data

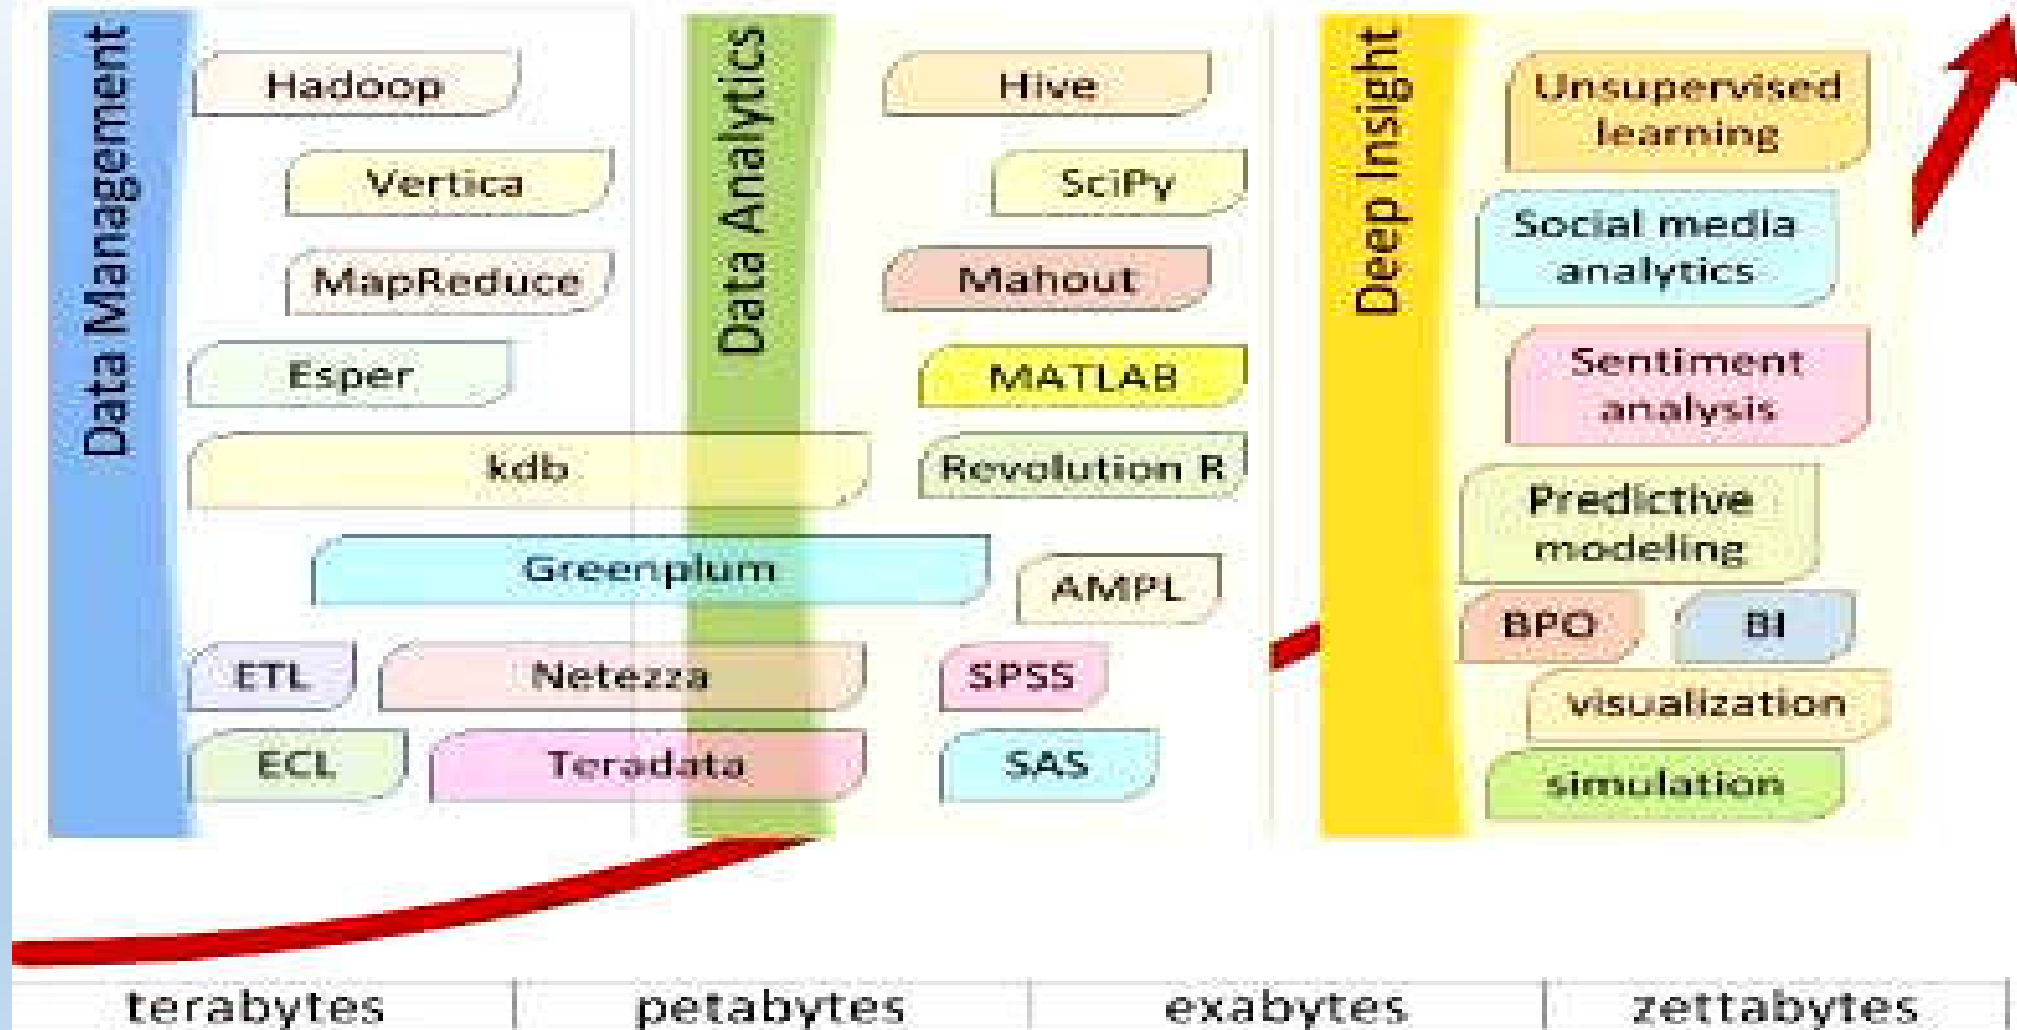
Focus on	Descriptive analytics Diagnostic analysis	Predictive analytics
Data Set	Limited dataset Cleaned data simple models	Large scale datasets More Types of data Raw Data Complex data Models
Supports	Causation: what happened and why?	Correlation: New Insights and more accurate answers.



# Enterprise data architecture using combined EDW/ Big data environment



# Big Data



# Technology for Big data analytics

## Document databases:









Pair each key with a complex data structure known as a document.

## Graph stores :

Used to store information about networks of data, such as social connections.

**Key-value stores** Every single item in the database is stored as an attribute name (or 'key'), together with its value.

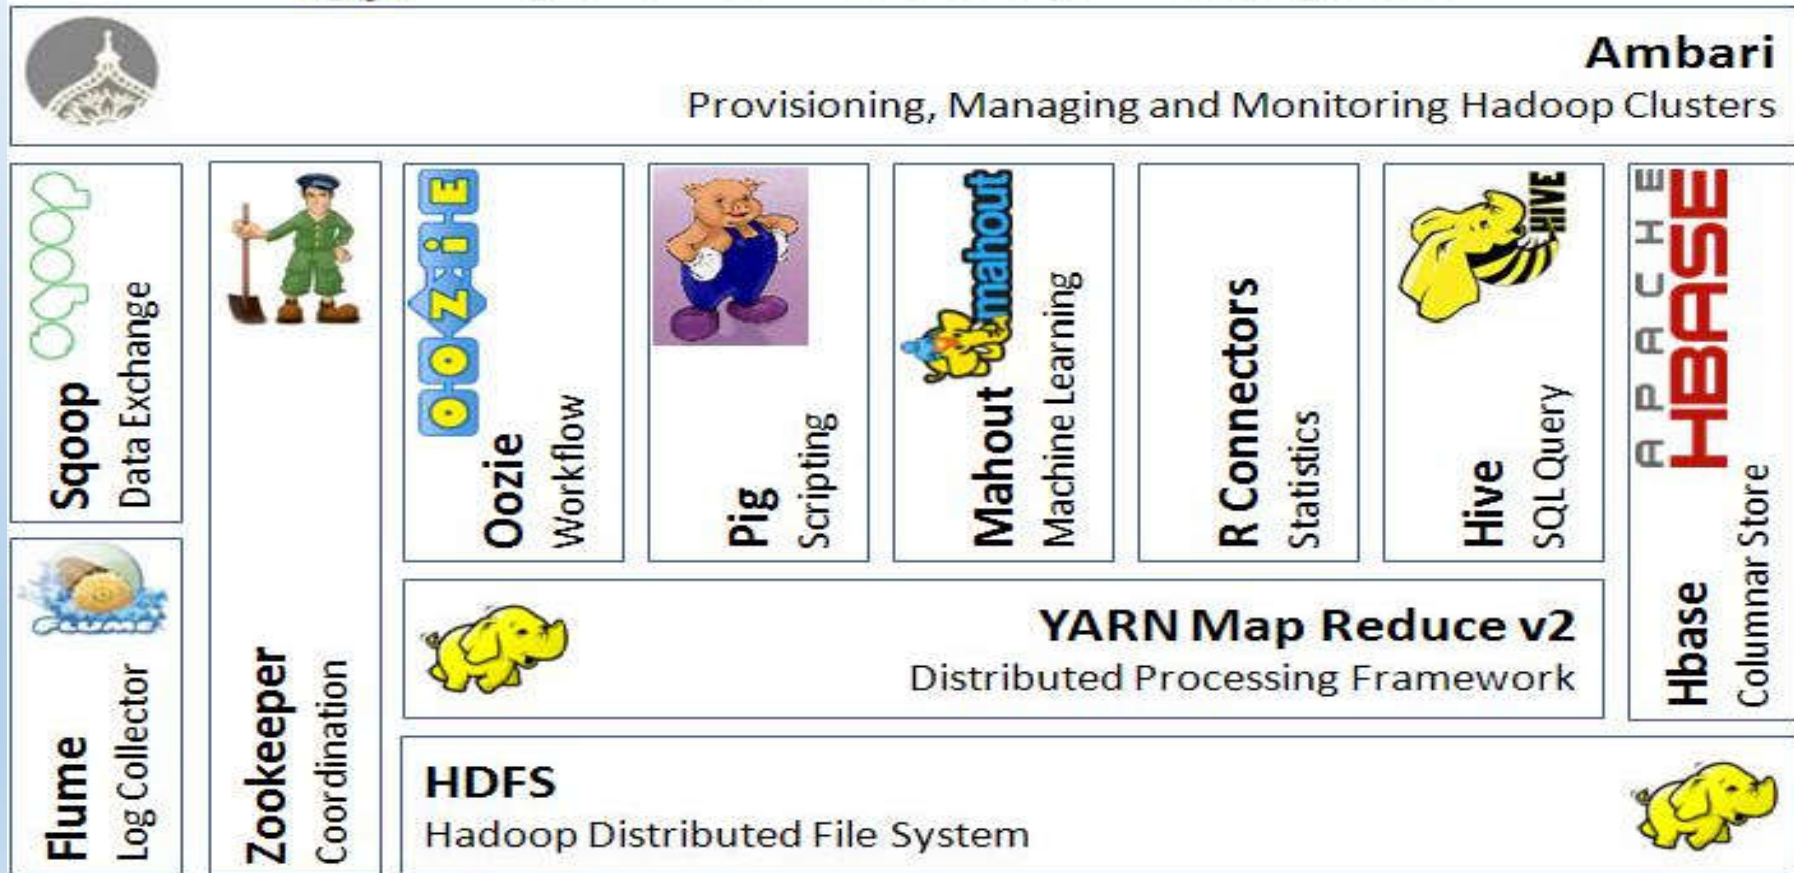
**Wide-column stores :** Store columns of data together, instead of rows

Type	Example	
Key-Value Store	 redis	 riak
Wide Column Store	 HBASE	 cassandra
Document Store	 mongoDB	 CouchDB relax
Graph Store	 Neo4j	 InfiniteGraph The Distributed Graph Database

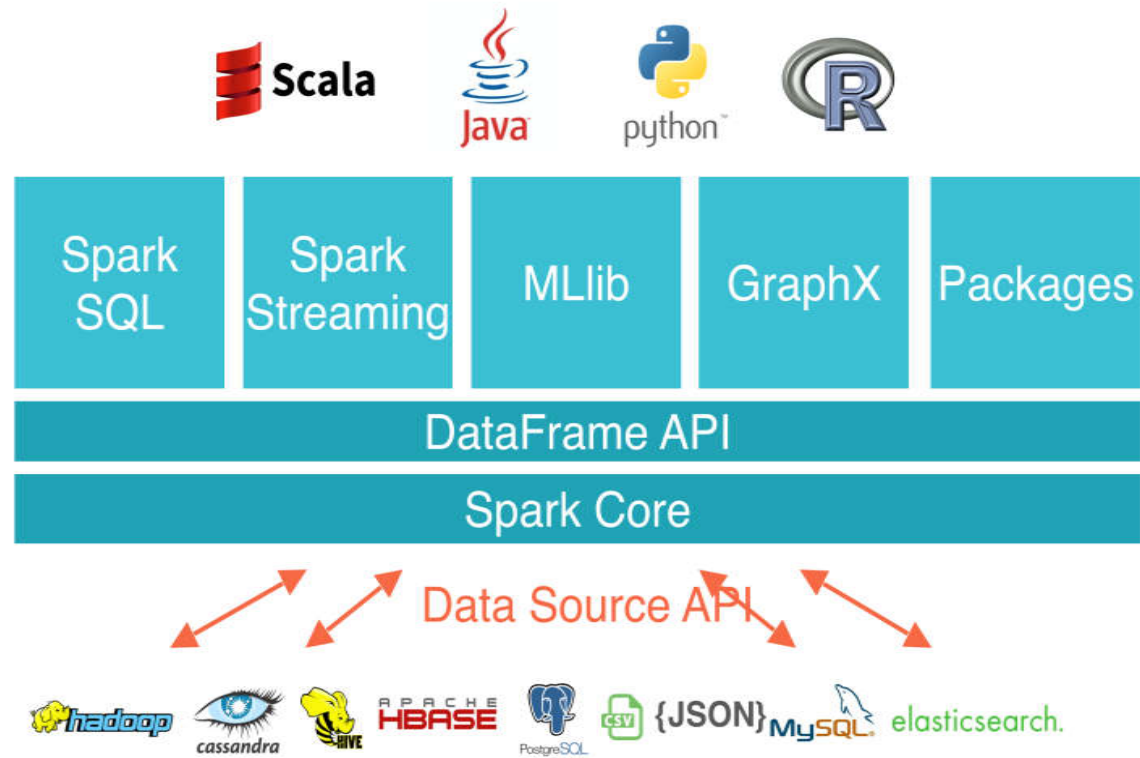
# Hadoop Eco systems



## Apache Hadoop Ecosystem



# Spark Ecosystem



## Top challenges facing Big data

- Scalability: Horizontal and vertical scalability
- Security: NOSQL dB
- Schema: Dynamic schemas
- Continuous availability: 24\*7
- Consistency: Opt for consistency
- Partition tolerant: Tolerance to both hardware and software failures
- Data quality: Data accuracy , completeness

## Kind of technologies looking toward to meet big data challenges

- Cheap and abundant storage
- Faster processors
- Affordable open sources, distributed platforms
- Parallel processing, clustering, Large grid environments
- Cloud computing and other flexible resource allocation arrangements.

## Skill set

DBMS and NOSQL Databases

Programming languages

Open source tools

Data warehousing / Data mining

Statistics analysis

Visualization

AI and ML



## Terminologies Used in Big data Environments

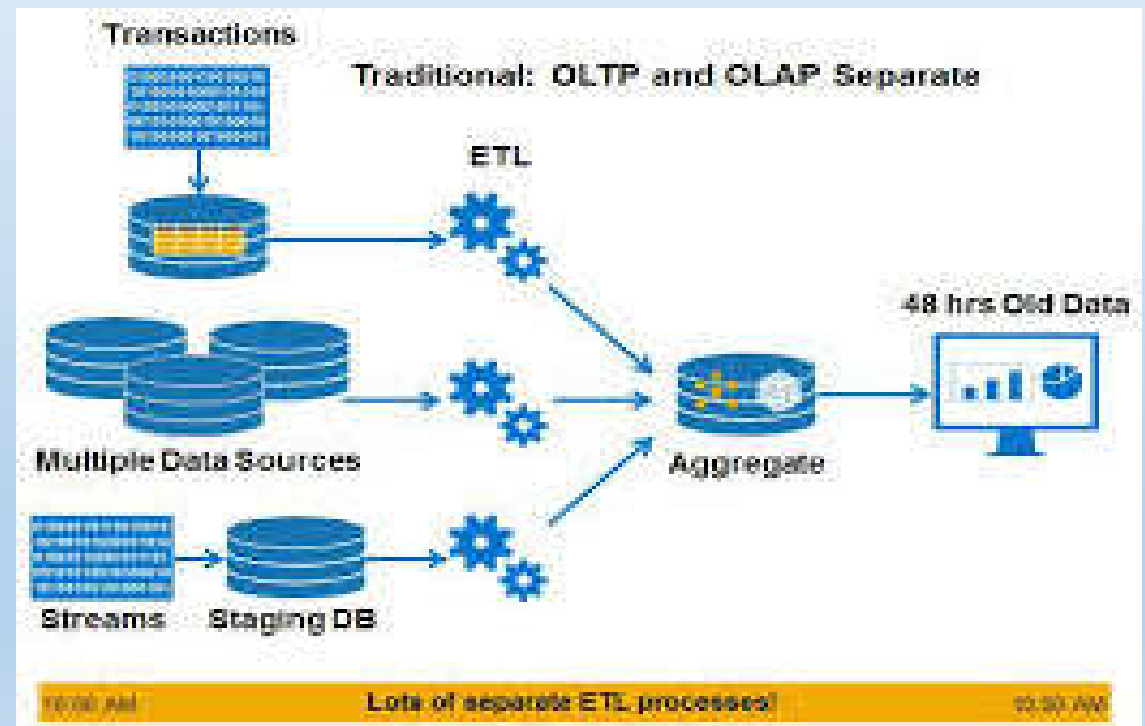
1. In-Memory Analytics
2. In-Database Processing
3. Massively Parallel Processing
4. Parallel System
5. Distributed System
6. Shared Nothing Architecture

## In- Memory analytics

- Data access from non- volatile storage such as hard disk and processed in RAM
- Shortened query response times
- Preprocess data and store it as cubes, tables, query sets etc, So it handle small set of records.
- All relevant data stored in RAM

## In- Database Processing


- It is also called as in- database analytics
- It works by fusing from - data warehouse / databases



## Symmetric Multiprocessor System (SMP)

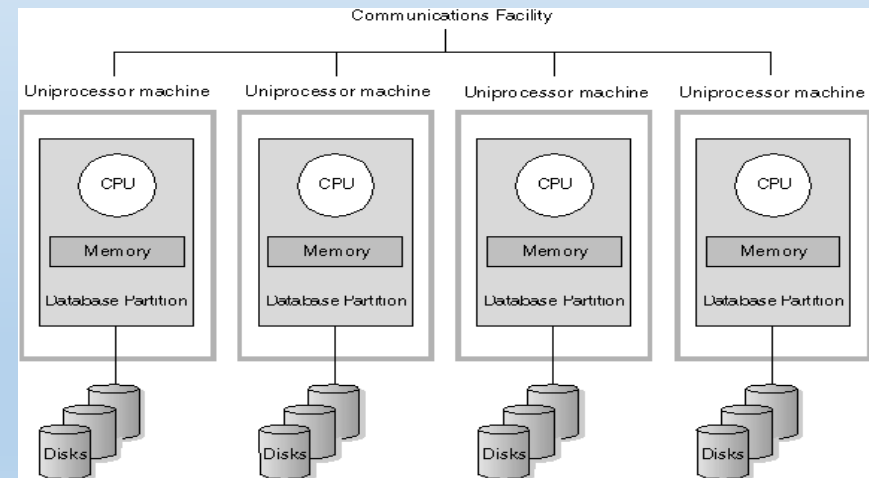
A **symmetric multiprocessor system (SMP)** is a **multiprocessor system** with centralized shared memory called **main memory (MM)** operating under a **single operating system** with two or more homogeneous processors

SMP is also called tightly coupled multiprocessing



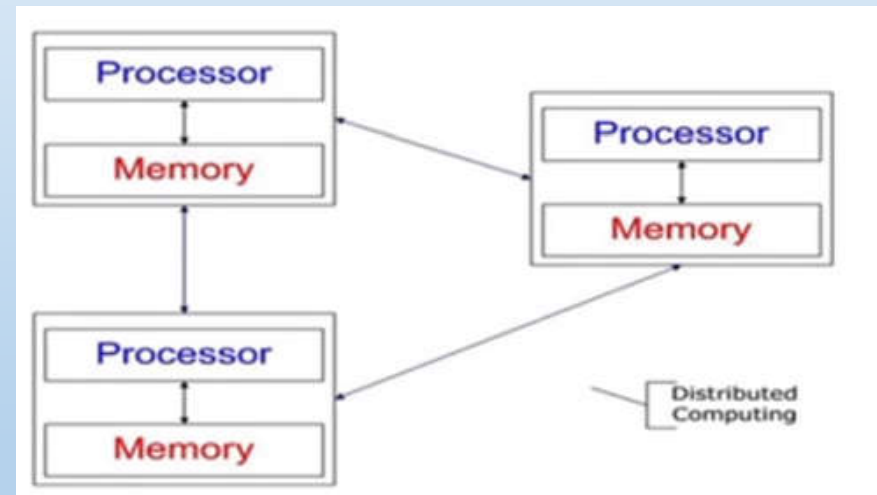
# Massively Parallel Processing

- MPP (**massively parallel processing**) is the coordinated **processing** of a program by multiple processors that work on different parts of the program, with each processor using its own operating system and memory .
- Each process works on different part of the program / data.
- Typically, MPP processors communicate using some messaging interface.
- Data and processing is governed solely by **performance considerations**.
- Used for scientific applications.
- May be
  - tightly (Shared memory)
  - loosely coupled ( no shared memory)



# Distributed Computing

- Different autonomous systems that have their own computational capacity and local memory.
- Systems Communicate with each other by message passing.
- Fault tolerant and high availability.
- Resource Sharing capability.
- Data and processing is governed solely by **application considerations**.
- Loosely Coupled System




# Multiprocessing Architecture

- Shared memory
- Shared disk
- Shared Nothing Architecture

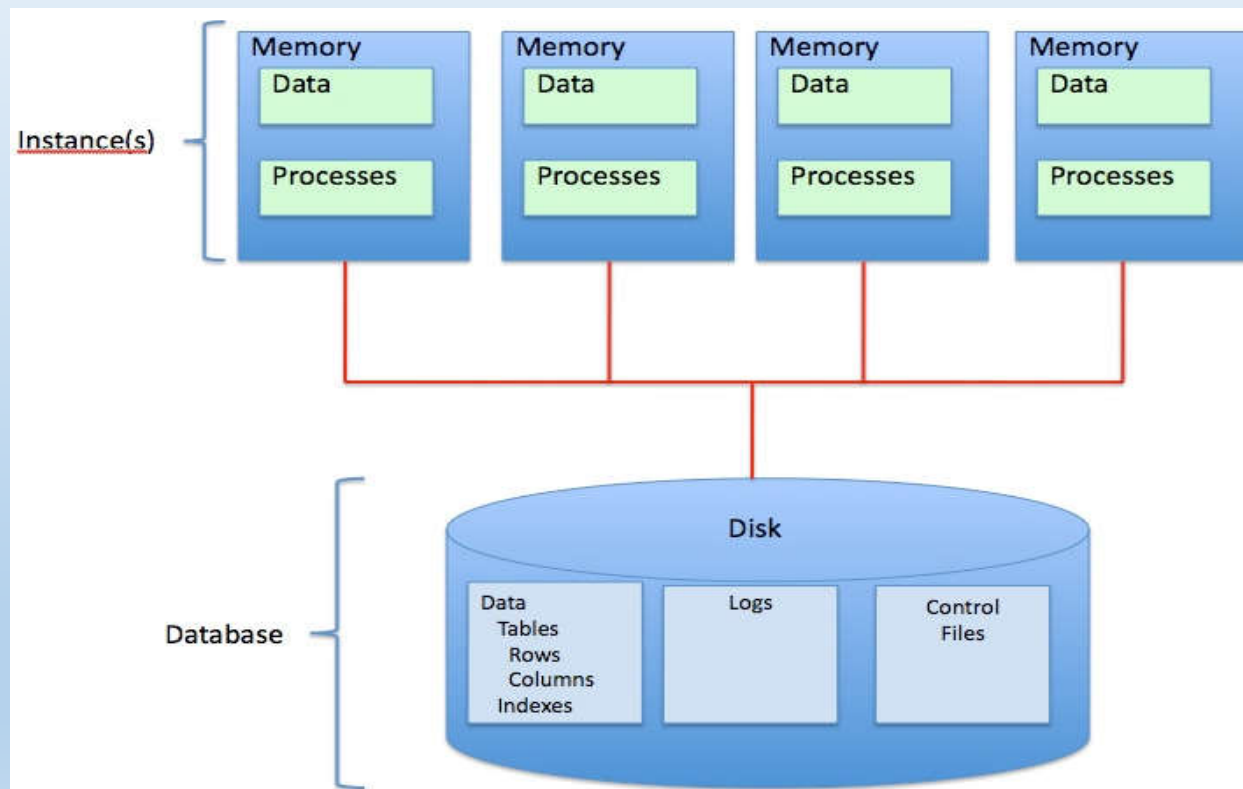
## shared memory

- **Shared memory** is **memory** that may be simultaneously accessed by multiple programs with an intent to provide communication among them or avoid redundant copies.
- **Shared memory** is an efficient means of passing data between programs.



# Shared Disk


A **shared disk** architecture (SD) is a distributed computing architecture in which all **disks** are accessible from all cluster nodes



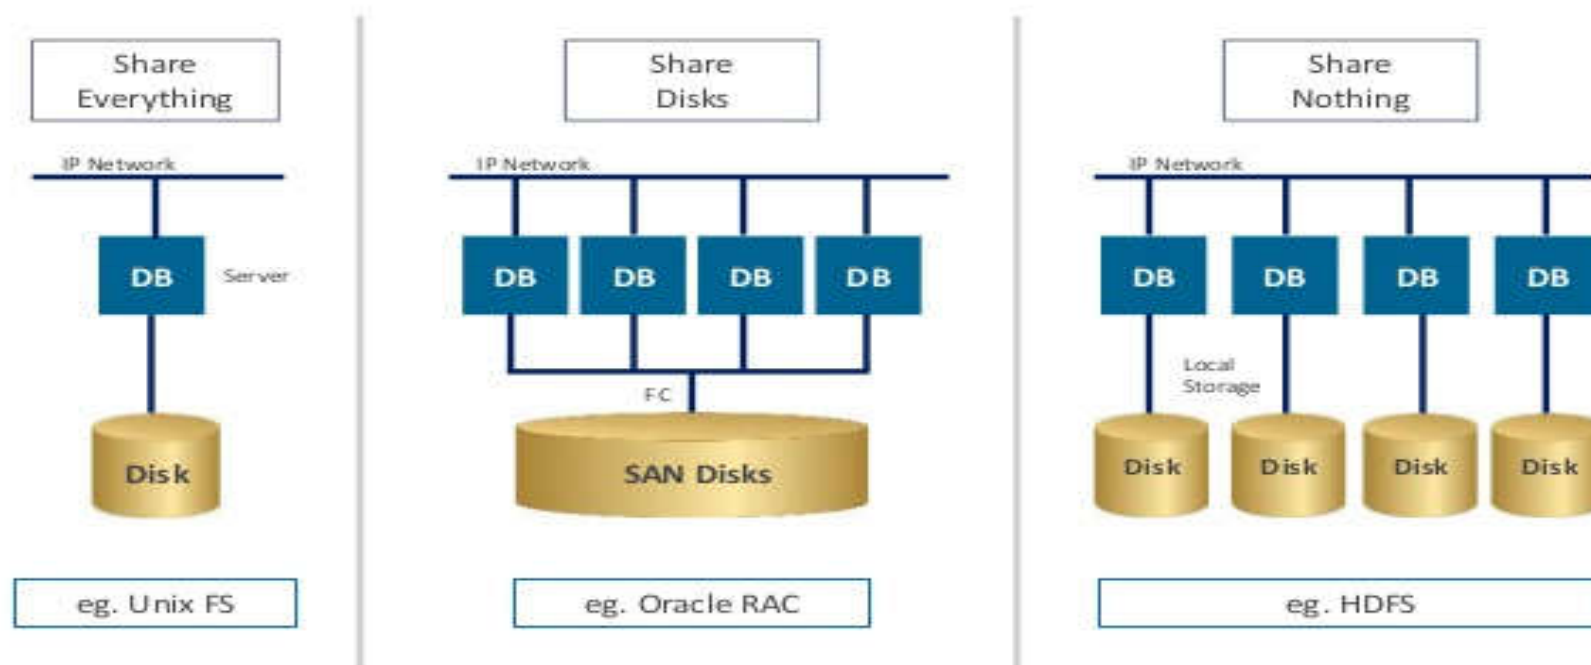


# Shared Nothing Architecture

- A **shared nothing architecture** (SN) is a distributed computing **architecture**
- Each node is independent and self-sufficient, and there is no single point of contention across the system.
- More specifically, none of the nodes share memory or disk storage.



## SHARE NOTHING ARCHITECTURE

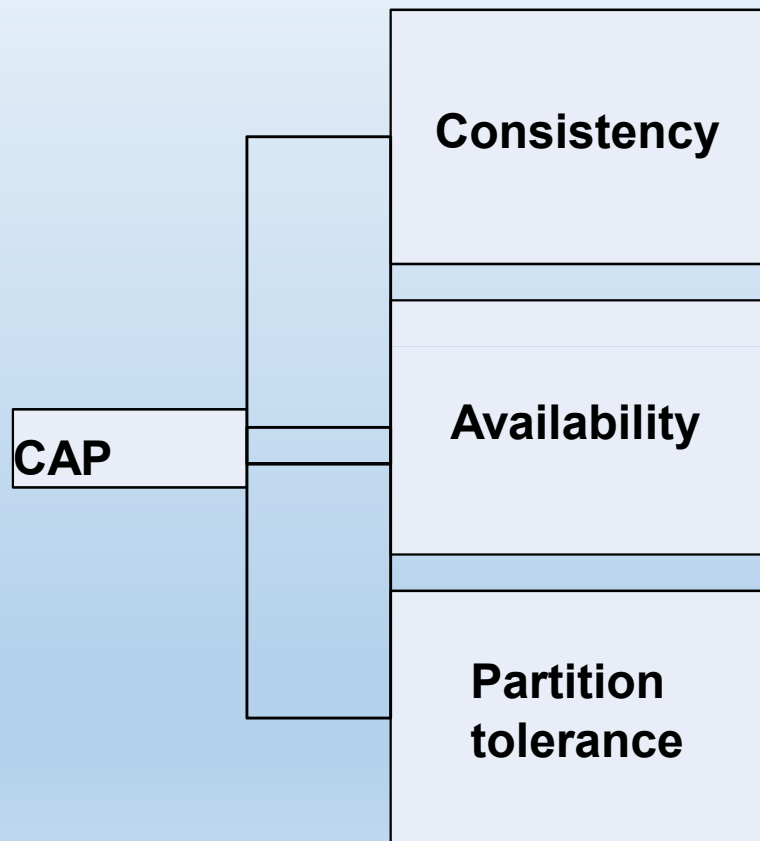


## **Advantages of shared Nothing architecture**

1. Fault Tolerance
2. Scalability

# CAP Theorem

## Brewer's CAP



## CAP Theorem


- **Consistency:**  
All nodes should see the same data at the same time
- **Availability:**  
Node failures do not prevent survivors from continuing to operate  
This condition states that **every request gets a response on success/failure of nodes.**  
Every client gets a response, regardless of the state of any individual node in the system.
- **Partition-tolerance:**  
The system **continues to operate despite network partitions failures.**  
partition-tolerant can sustain any amount of network failure that doesn't result in a failure of the entire network.  
Data records are **sufficiently replicated across combinations of nodes** and networks to keep the system up through intermittent outages.

A distributed system can satisfy any two of these guarantees at the same time **but not all three**

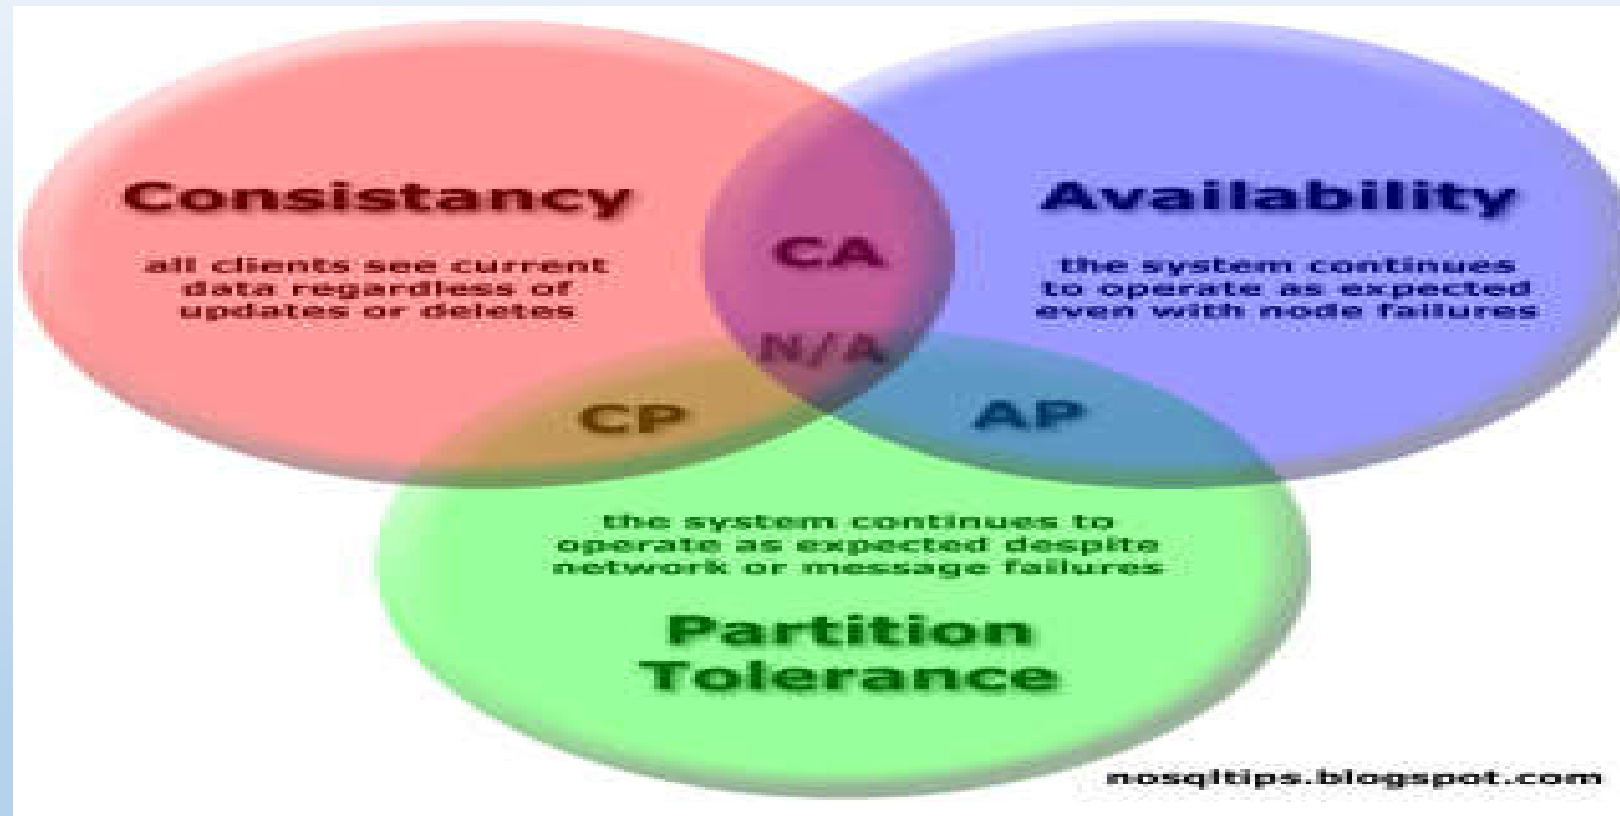
In an **available but not partition-tolerant system**, Y would be allowed to process its request because it can reach X to obtain the lock. Z's request would be blocked because X is unreachable.

In a **partition-tolerant but not available system**, Z would be allowed to process its request because it is part of the quorum group (X's lock will be broken). Y's request would be blocked because it is not part of the quorum group.

In a system that is **both available and partition-tolerant**, both requests would be allowed to progress. Y would return current data as possibly modified by X, while Z would return possibly stale data. Stale data could possibly mean no data in cases where there is no replica available with Quorum nodes. Consistency is obviously sacrificed in this case.




## CAP theorem






# CAP Theorem


Consistency



Availability



Partition tolerance



## Why this is important?

- The future of databases is **distributed** (Big Data Trend, etc.)
- CAP theorem describes the **trade-offs** involved in distributed systems
- A proper understanding of CAP theorem is essential to **making decisions** about the future of distributed database **design**
- Misunderstanding can lead to **erroneous or inappropriate** design choices

Fox&Brewer “CAP Theorem”:  
C-A-P: choose two.

CA: available, and consistent, unless there is a partition.

Claim: every distributed system is on one side of the triangle.

CP: always consistent, even in a partition, but a reachable replica may deny service without agreement of the others (e.g., quorum).

**Availability**

A

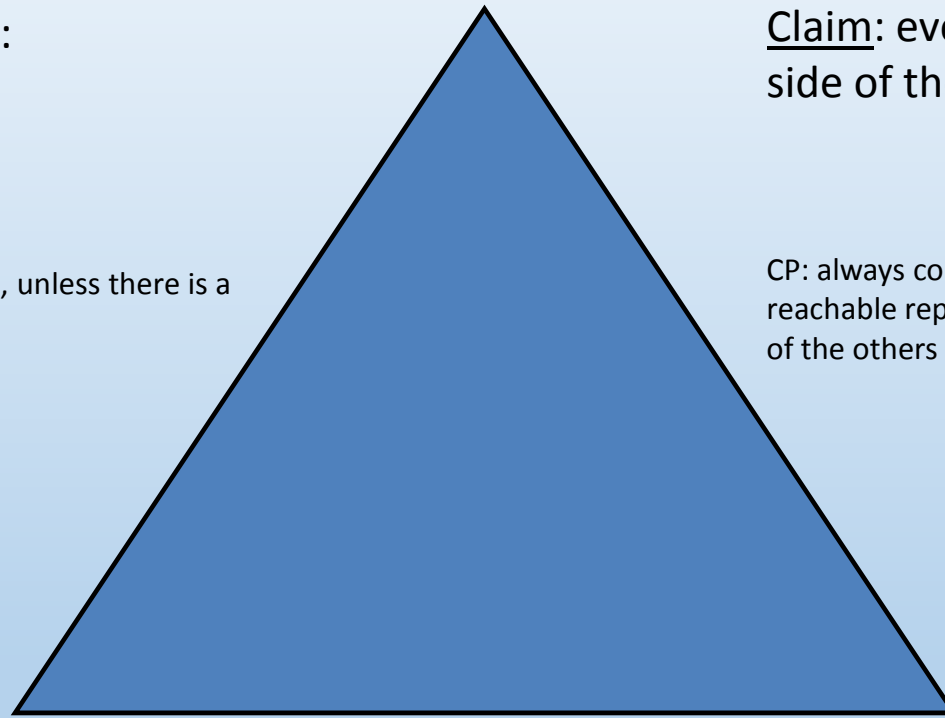
AP: a reachable replica provides service even in a partition, but may be inconsistent.

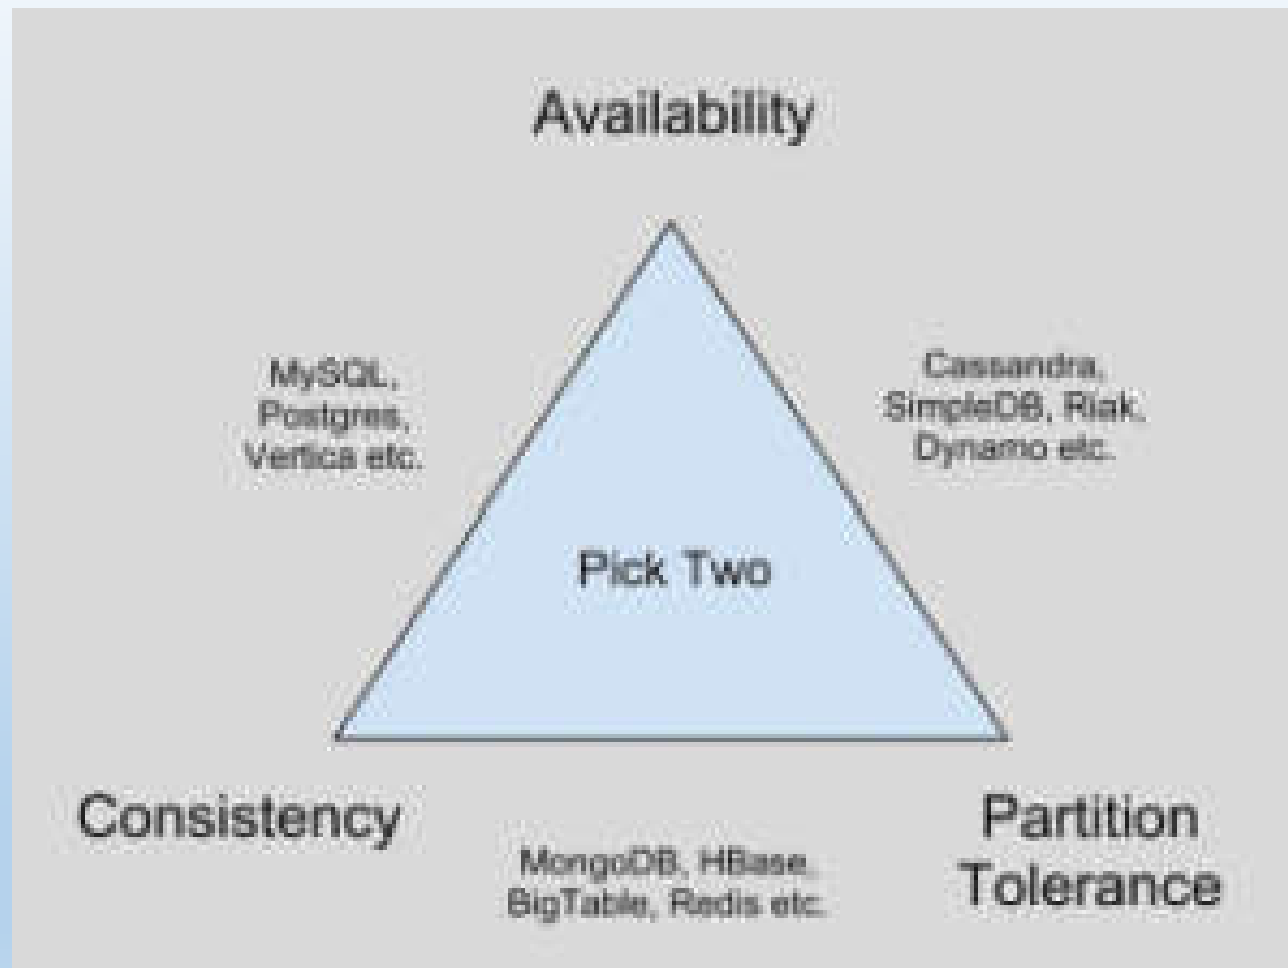
Partition-resilience

P

consistency

C





## When to consider consistency over availability and Vice -versa

1. Choose availability over consistency when your business requirements allow some flexibility around when data in the system synchronizes.
2. Choose consistency over availability when your business requirements demand atomic reads and writes.

## Types of Consistency

### Strong Consistency

After the update completes, **any subsequent access** will return the **same** updated value.

### Weak Consistency

It is **not guaranteed** that subsequent accesses will return the updated value.

### Eventual Consistency

Specific form of weak consistency

It is guaranteed that if **no new updates** are made to object, **eventually** all accesses will return the last updated value (e.g., *propagate updates to replicas in a lazy fashion*)

## Eventual Consistency Variations

### Causal consistency

Processes that have causal relationship will see consistent data

### Read-your-write consistency

A process always accesses the data item after it's update operation and never sees an older value

### Session consistency

As long as session exists, system guarantees read-your-write consistency  
Guarantees do not overlap sessions

## Eventual Consistency Variations

### Monotonic read consistency

If a process has seen a particular value of data item, any subsequent processes will never return any previous values

### Monotonic write consistency

The system guarantees to serialize the writes by the *same* process

### In practice

A number of these properties can be combined

Monotonic reads and read-your-writes are most desirable



## Eventual Consistency

### - A Facebook Example

Bob finds an interesting story and shares with Alice by posting on her Facebook wall


Bob asks Alice to check it out

Alice logs in her account, checks her Facebook wall but finds:

- **Nothing is there!**




?



## Eventual Consistency - A Facebook Example

Bob tells Alice to wait a bit and check out later  
Alice waits for a minute or so and checks back:  
- **She finds the story Bob shared with her!**



## Eventual Consistency - A Facebook Example

Reason: it is possible because Facebook uses an **eventual consistent model**

Why Facebook chooses eventual consistent model over the strong consistent one?

- Facebook has more than 1 billion active users
- It is non-trivial to efficiently and reliably store the huge amount of data generated at any given time
- Eventual consistent model offers the option to **reduce the load and improve availability**

# BASE

- Basically Available Soft state Eventual consistency
- Where it is used? – Distributed Computing
- Why?- To achieve high availability
- How it is achieved? - No new updates made to the data for a stipulated period of time , eventually all accesses to this data will return the updated value.
- What is replica convergence? system that has achieved eventual consistency is said to have converged.
- Conflict resolution: solved by
  1. Read repair
  2. Write repair
  3. Asynchronous repair

## ACID vs BASE

### ACID

- Strong consistency for transactions highest priority
- Availability less important
- Complex mechanisms

### BASE

- Availability and scaling highest priorities
- Weak consistency
- Simple and fast

# Thank you